Soutenance finale

Théo Ripoll - Tom Genlis - Arnaud Baradat - Quentin Fisch

01

Dataset Presentation

Chosen dataset

02

EDA

Exploratory data analysis

03

Analysis - Network

Methods and processes

04

Analysis - Physical

Methods and processes

05

Demo

Key results & adversarial attack

06

Conclusion

Recap of results and takeaways

01 Dataset Presentation

Chosen dataset

Hardware In The Loop

Physical and network data from a Water Distribution Testbed, simulating water flow with real and virtual components to analyze effects of cyber and physical attacks in a 2-hour period.

Timeline of attacks

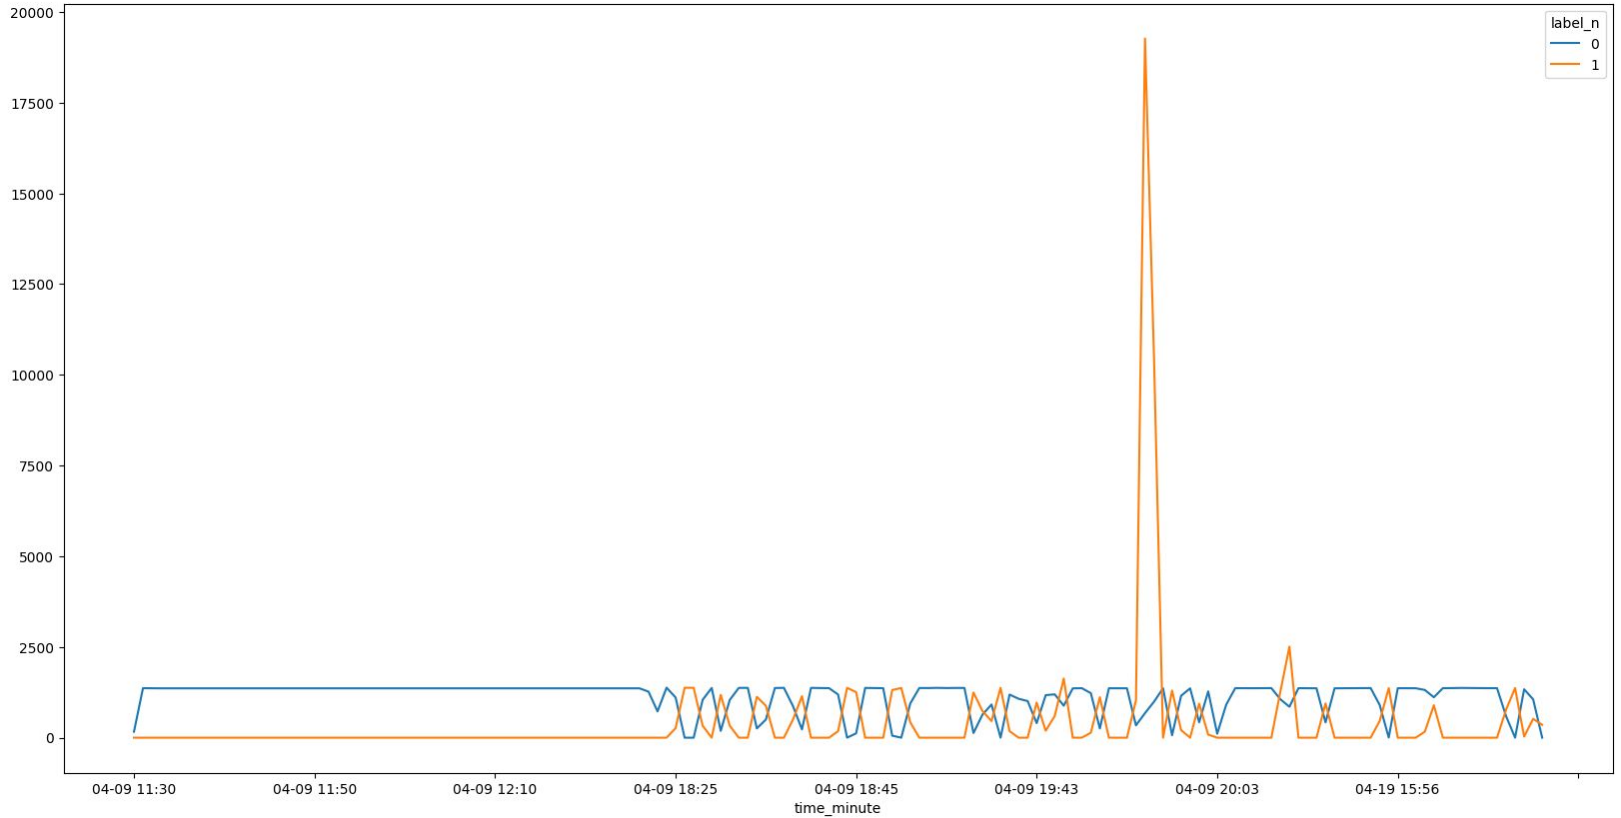

Timeline of attacks (2)

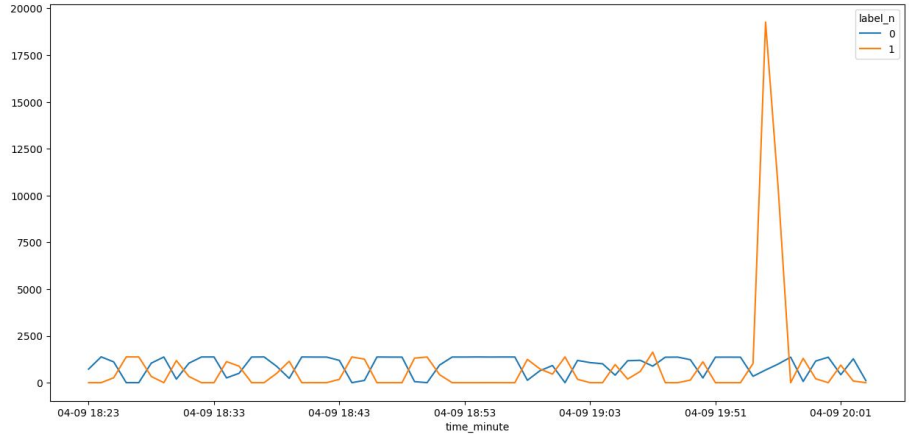

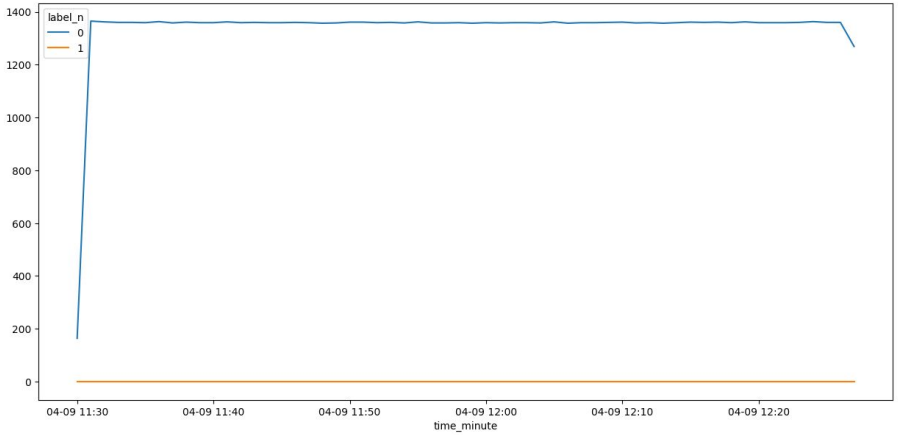

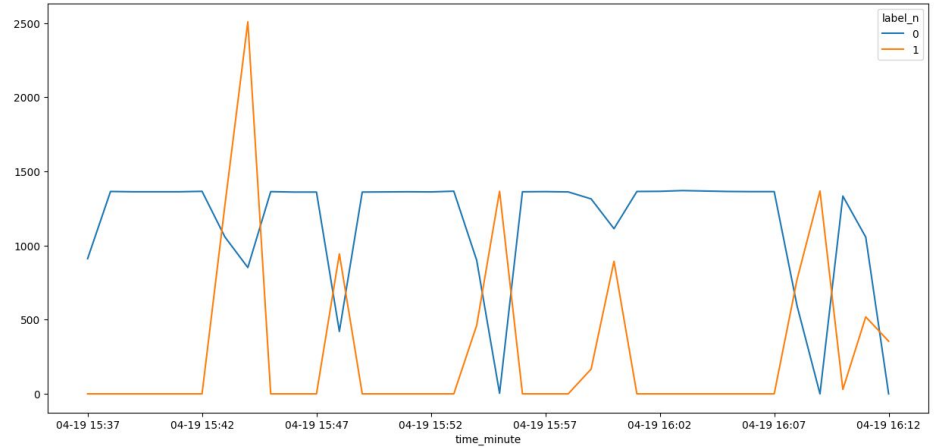

Physical sensors behavior

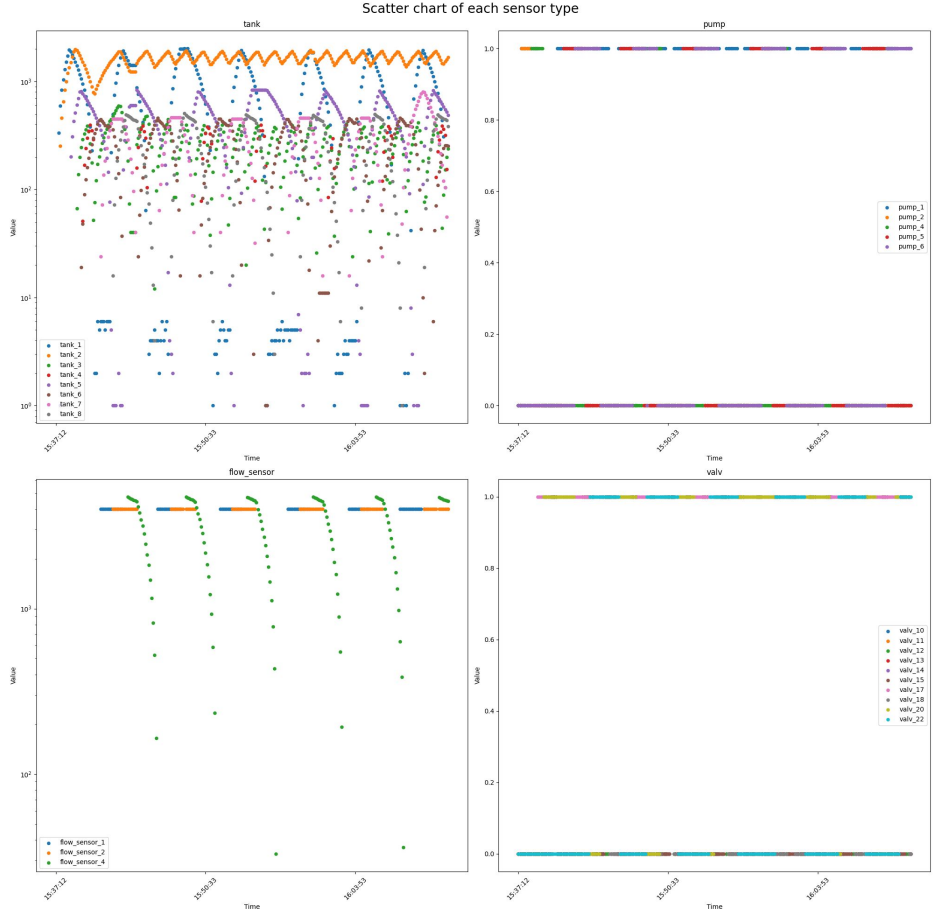

02

EDA

Exploratory Data Analysis

Labels - Binary class

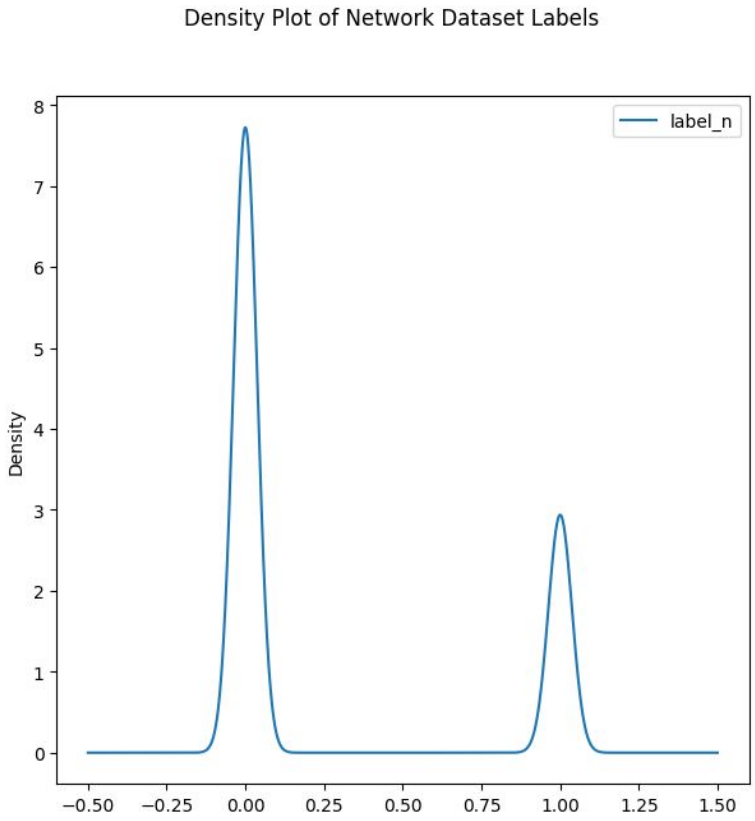

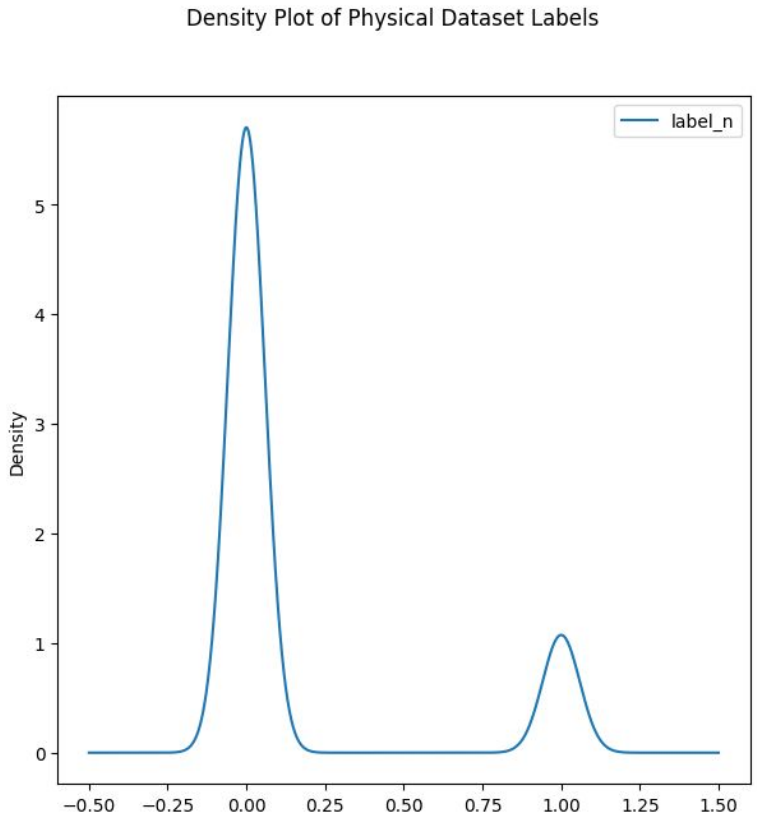

Labels - Multi-class

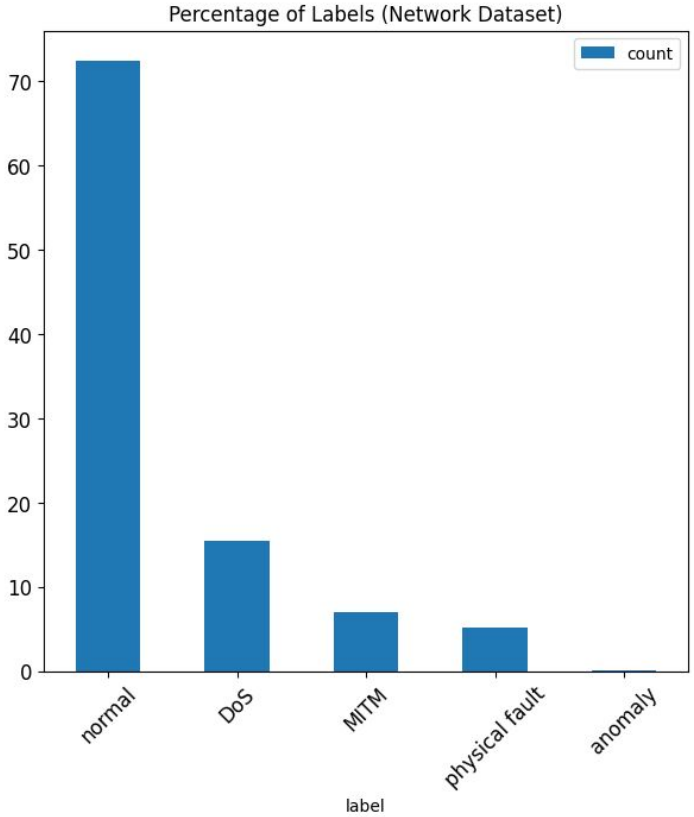

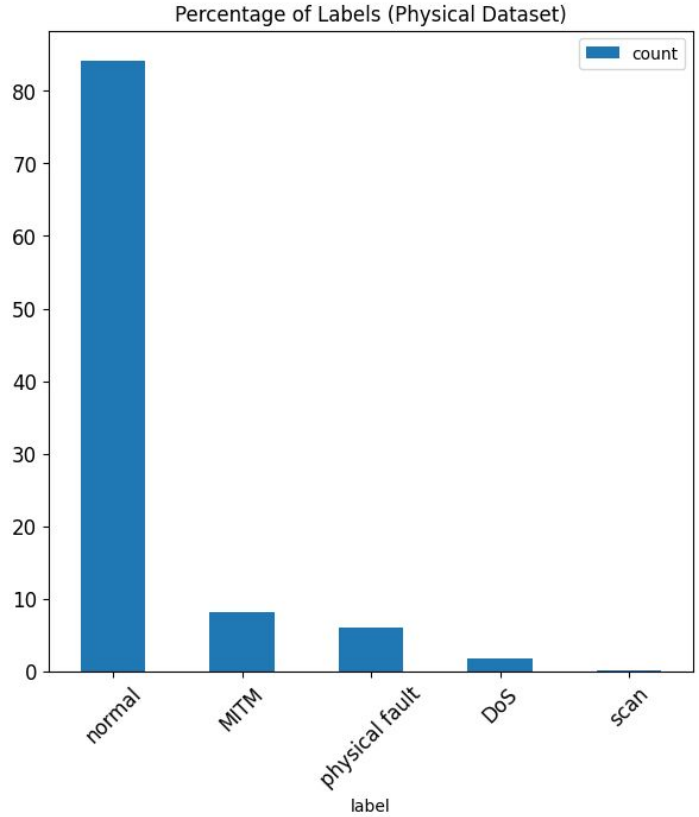

Quick rewind to our timeline

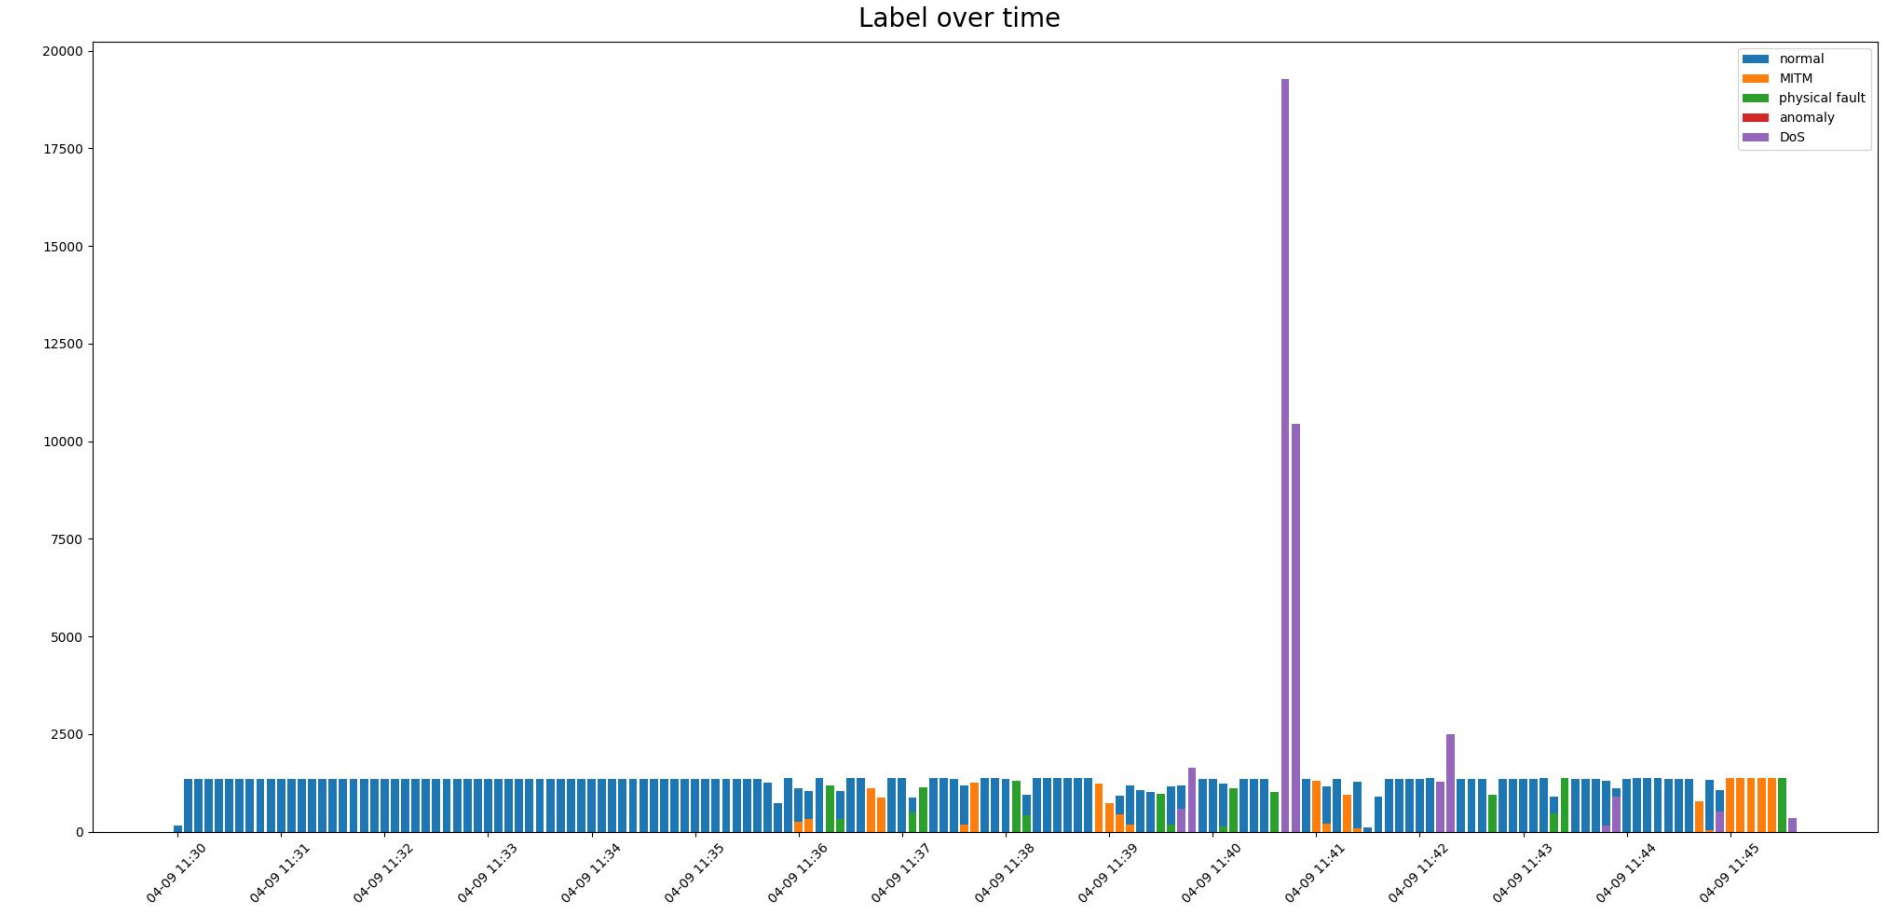

Network features

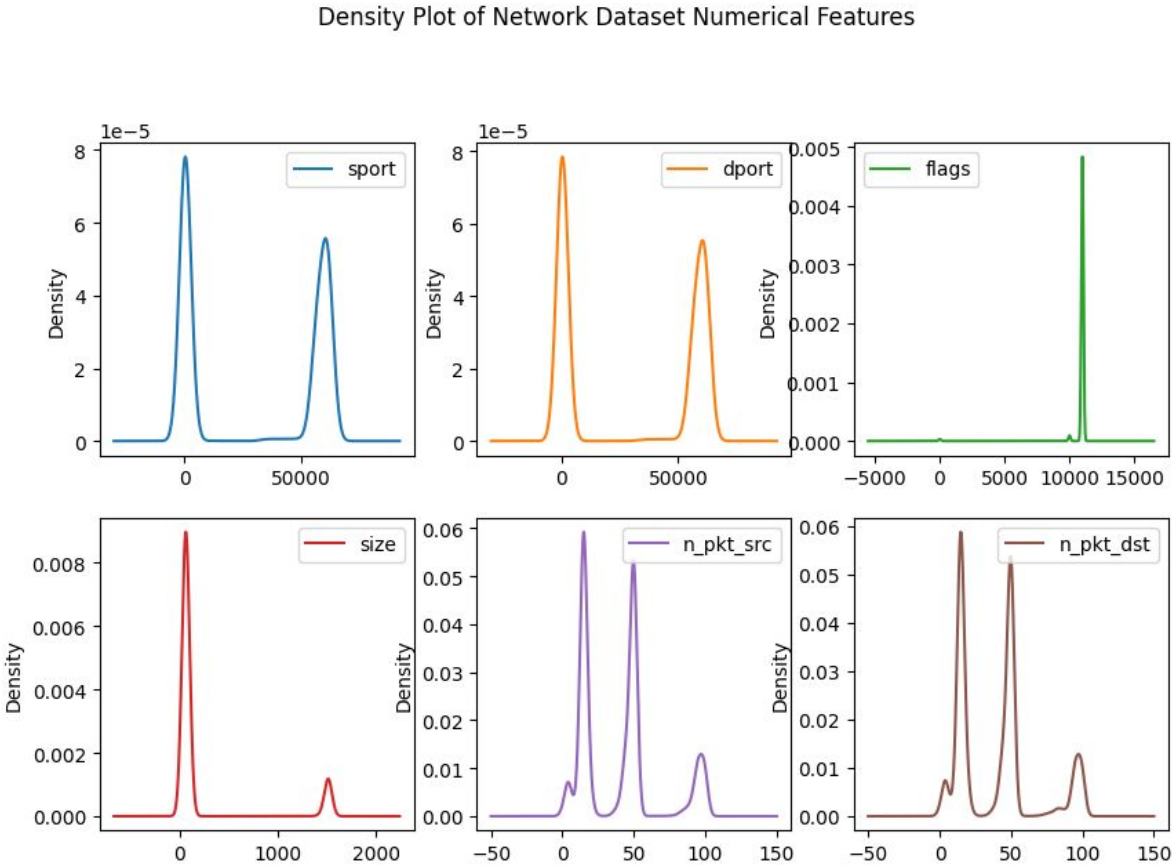

Network features (2)

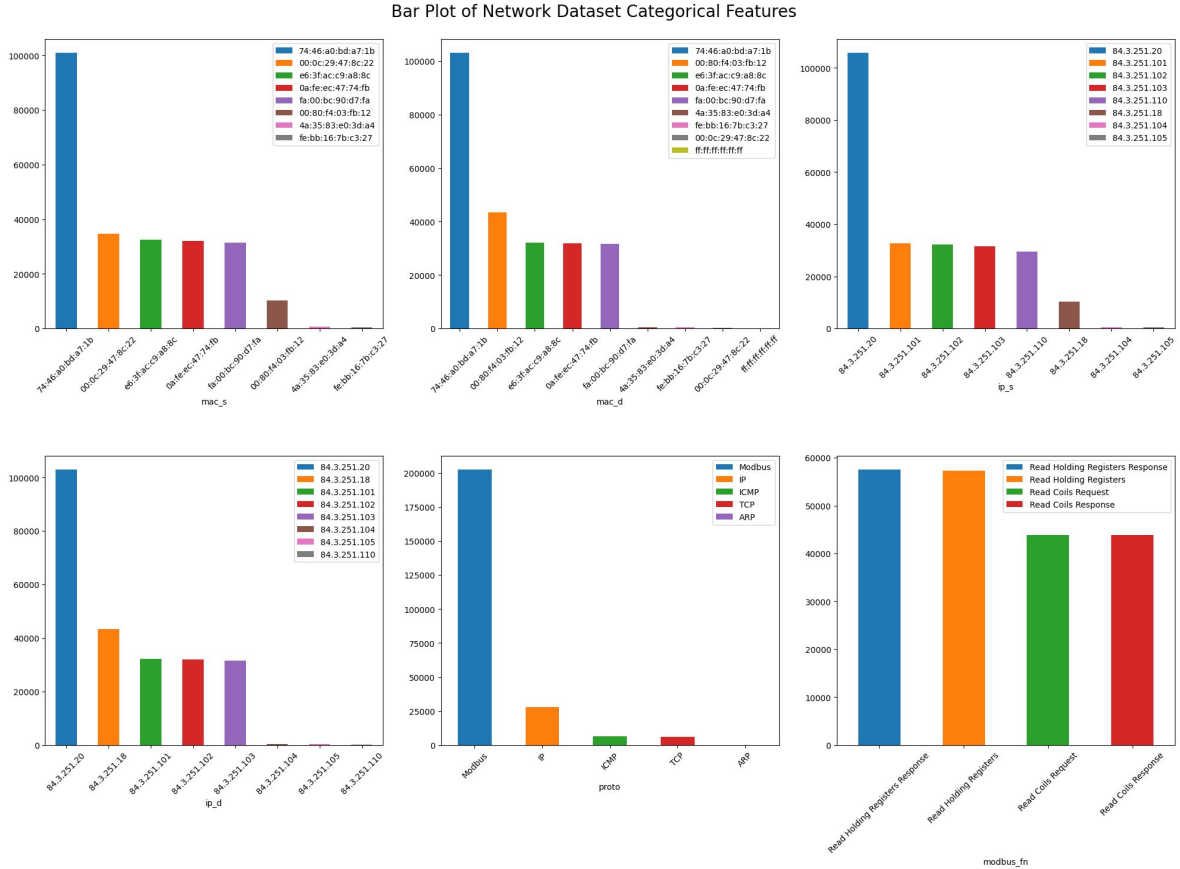

Network features (3)

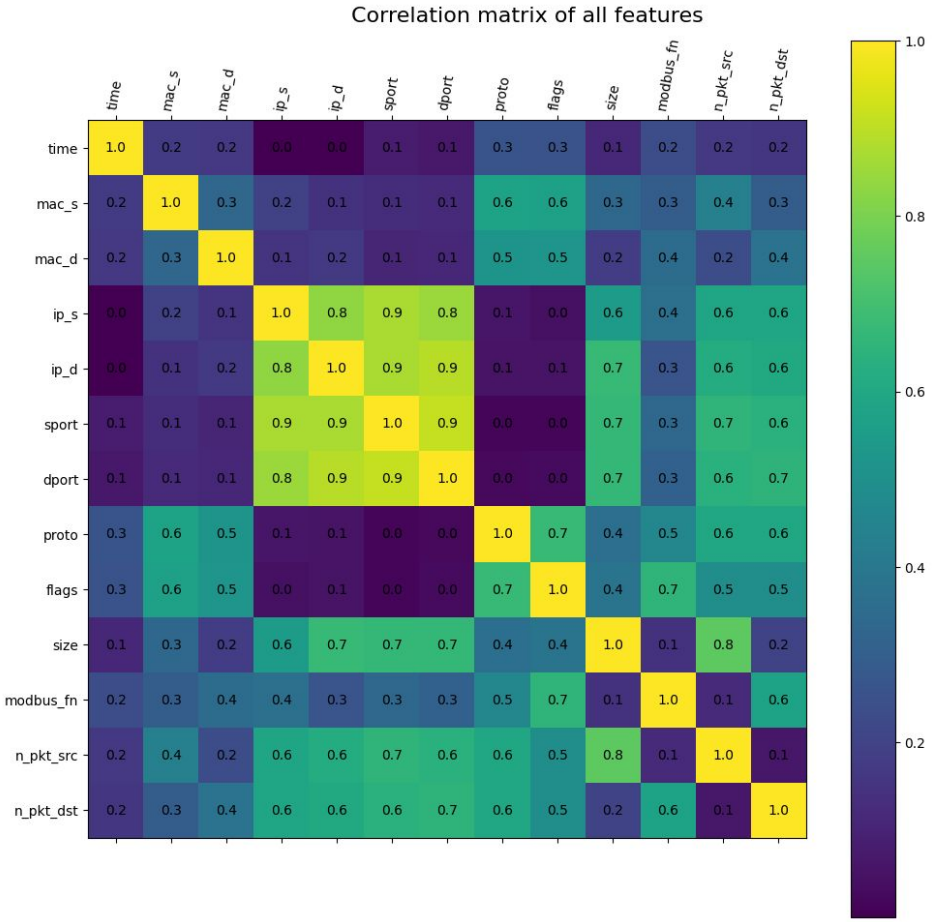

Physical Features

Physical Features (2)

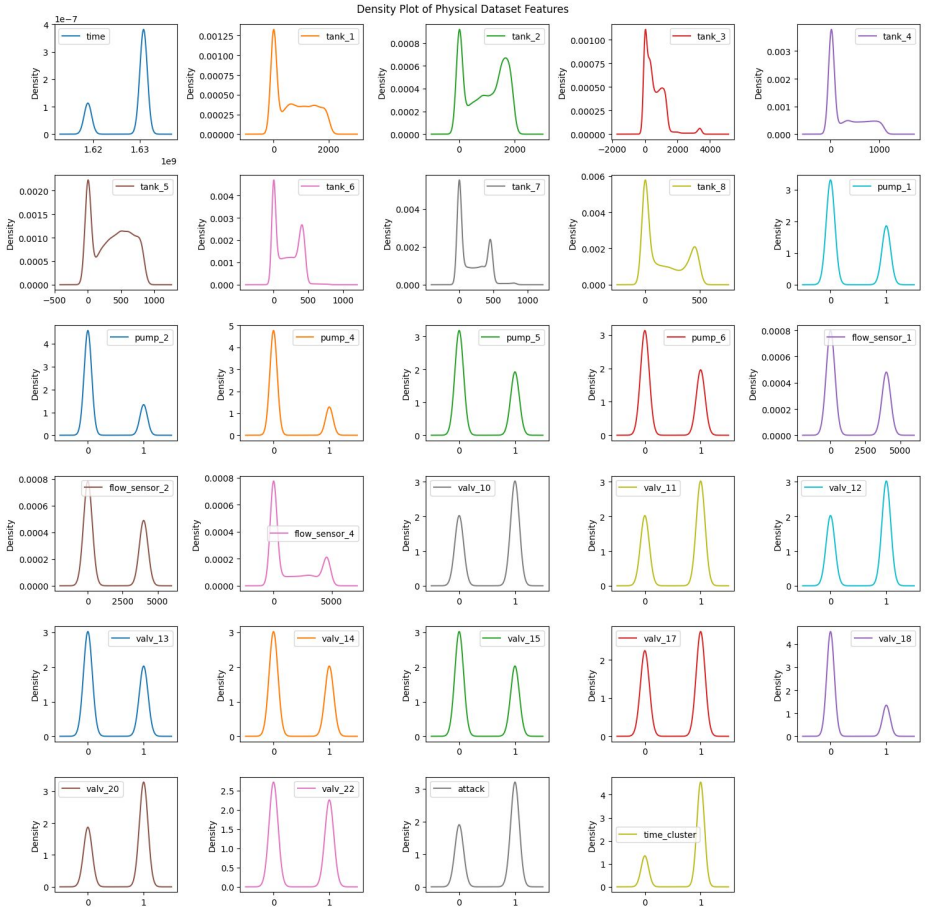

03 Analysis - Network

Methods and processes

Process

- Focus on multi-class classification
- Use a StandardScaler
- Separate analysis:
  - Using full contextual information
  - Without contextual information
- Use the following metrics:
  - Accuracy
  - Recall
  - F1
  - MCC
  - Balanced accuracy
- Plot and analyse feature importance
- Same for network and physical datasets

Non-supervised Algorithms

Isolation Forest

- Without a fixed contamination rate:
  - **12505** outliers detected (5%)
  - **7695** are real anomalies
  - **61.5%** of precision
- Contamination rate at 27.5%:
  - 66k anomalies to find
  - **37940** real anomalies detected
  - **57.7%** of precision

Local Outlier Factor

- Too slow for any production use
- Contamination rate at 27.5%:
  - **30919** outliers detected (12%)
  - Only **3599** are real anomalies

Deep Learning

Neural network

- Tried multiple architectures
- Final one has:
  - 3 hidden layers (1024, 256, 64)

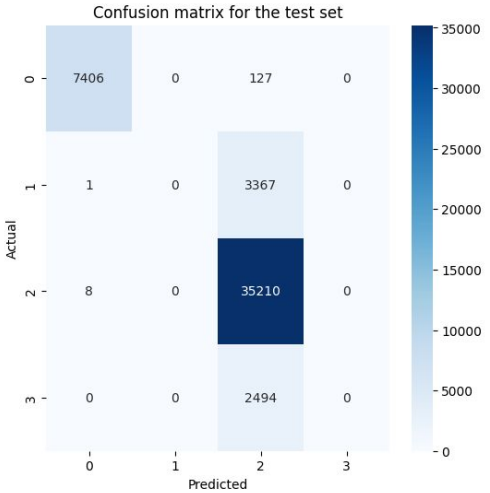

LSTM

- Similar results to the neural network

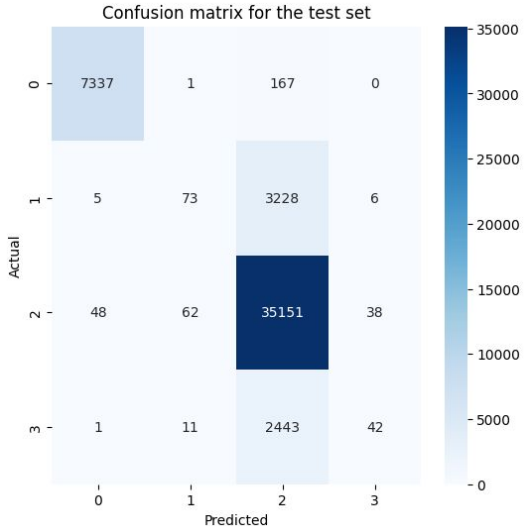

Classifiers - Decision Tree

With contextual information

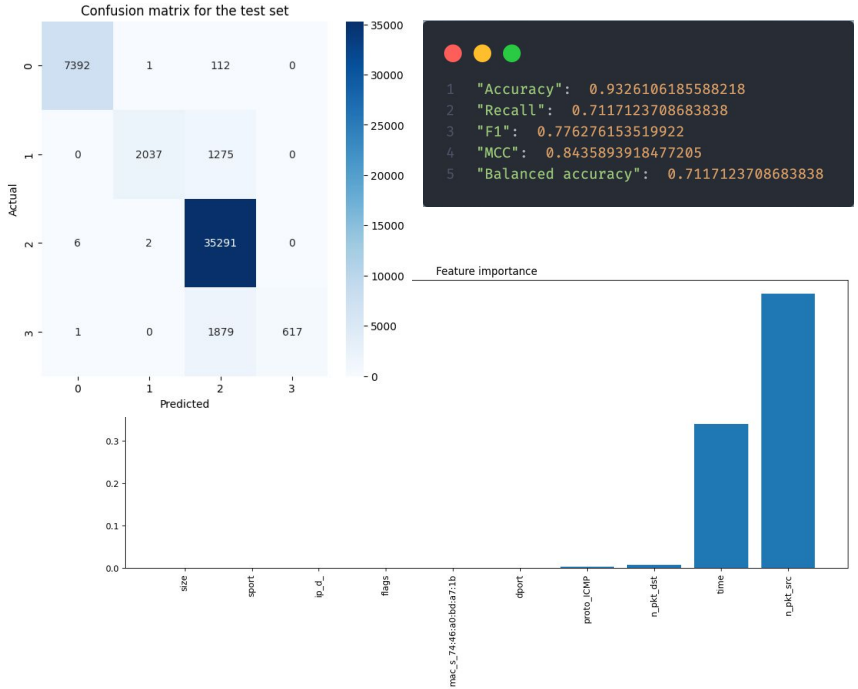


```
"Accuracy":  0.9326106185588218
"Recall":  0.7117123708683838
"F1":  0.776276153519922
"MCC":  0.8435893918477205
"Balanced accuracy":  0.7117123708683838
```

No contextual information

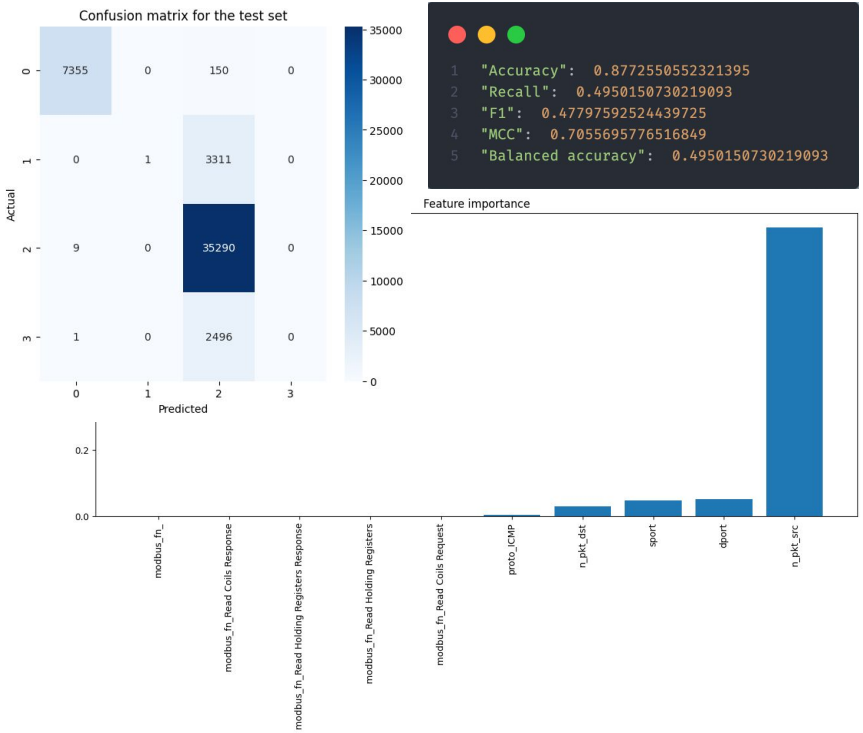


```
"Accuracy":  0.8772550552321395
"Recall":  0.4950150730219093
"F1":  0.47797592524439725
"MCC":  0.7055695776516849
"Balanced accuracy":  0.4950150730219093
```

Classifiers - Random Forest

With contextual information

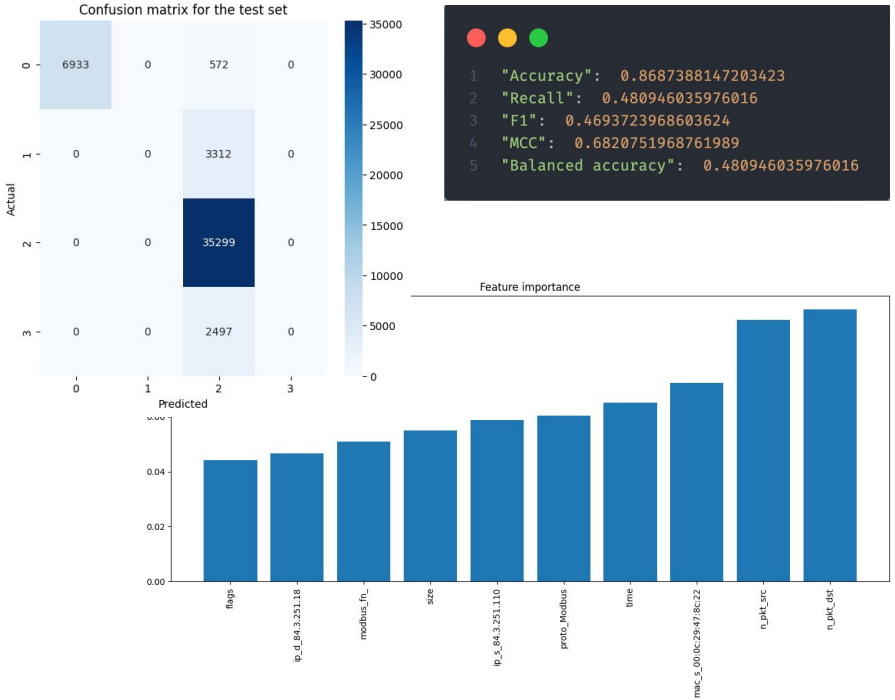


```
"Accuracy":  0.8687388147203423
"Recall":  0.480946035976016
"F1":  0.4693723968603624
"MCC":  0.6820751968761989
"Balanced accuracy":  0.480946035976016
```

No contextual information

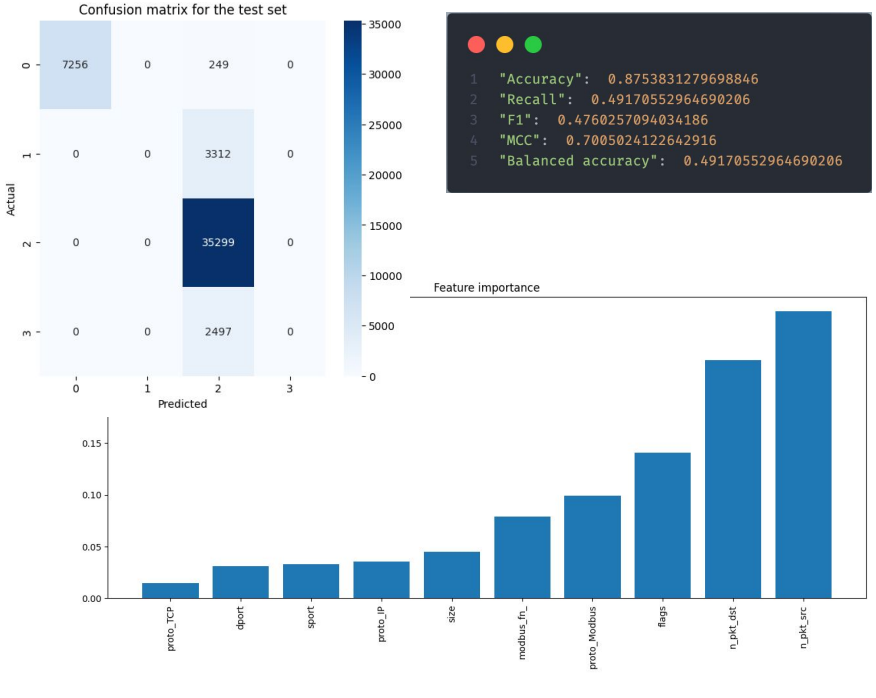


```
"Accuracy":  0.8753831279698846
"Recall":  0.49170552964690206
"F1":  0.4760257094034186
"MCC":  0.7005024122642916
"Balanced accuracy":  0.49170552964690206
```

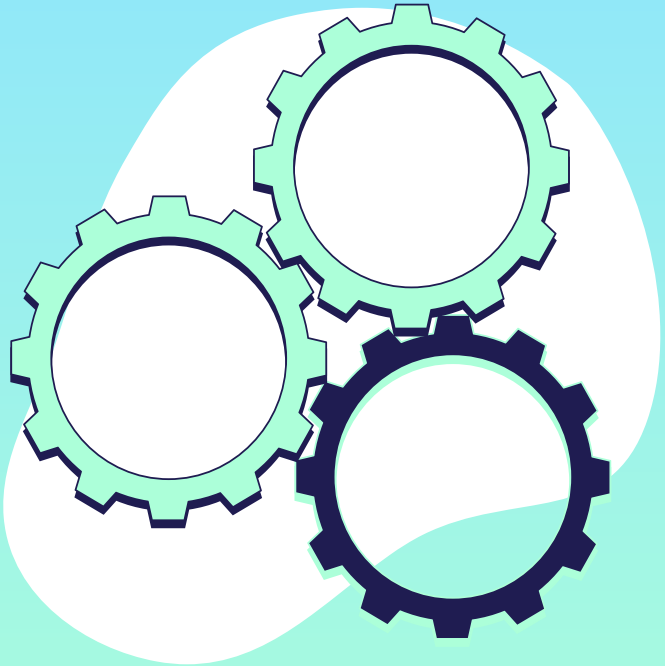

04 Analysis - Physical

Methods and processes

Non-supervised Algorithms

Isolation Forest

- Without a fixed contamination rate:
  - **6177** outliers detected
  - **1180** are real anomalies
  - **19%** of precision
- Contamination rate at 16%:
  - 66k anomalies to find
  - **37940** real anomalies detected
  - **57.7%** of precision

Local Outlier Factor

- Too slow for any production use
- **71** outliers detected (12%)
- Only **24** are real anomalies

Deep Learning

Neural network

- Tried multiple architectures
- Final one has:
  - 3 hidden layers (1024, 256, 64)

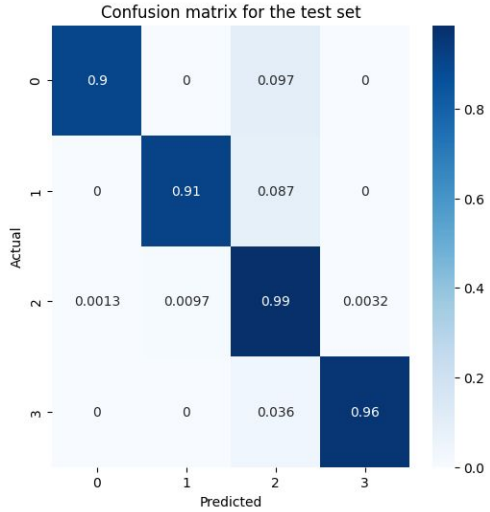

LSTM

- Similar results to the neural network

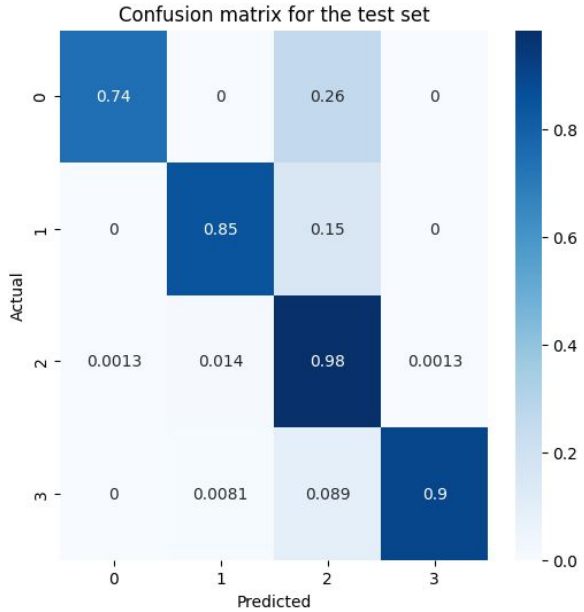

Classifiers - Decision Tree

With contextual information

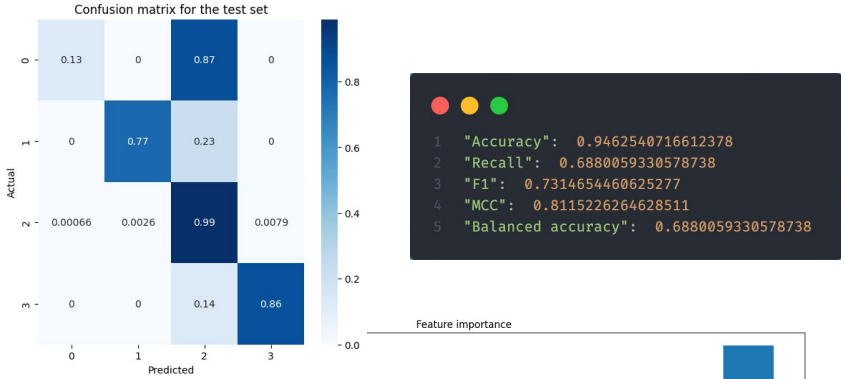


```
"Accuracy":  0.9462540716612378
"Recall":  0.6880059330578738
"F1":  0.7314654460625277
"MCC":  0.8115226264628511
"Balanced accuracy":  0.6880059330578738
```

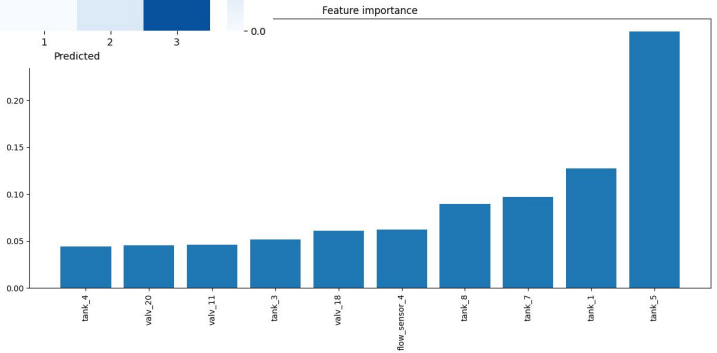

No contextual information and with balanced classes

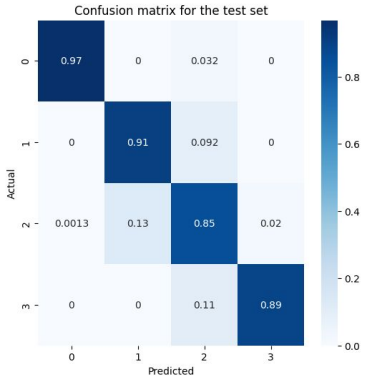


```
"Accuracy":  0.8604777415852335
"Recall":  0.9032609495946987
"F1":  0.8198354103357238
"MCC":  0.663623364956895
"Balanced accuracy":  0.9032609495946987
```

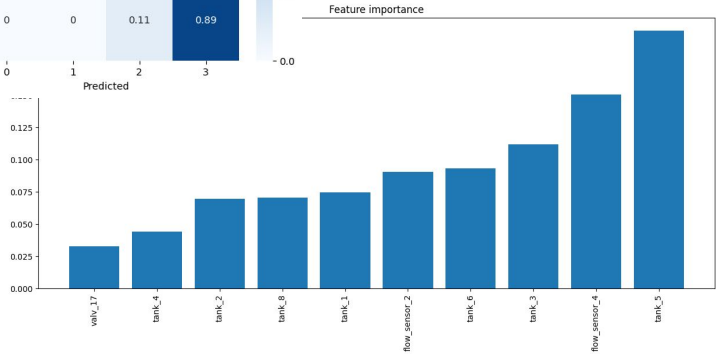

Classifiers - Random Forest

With contextual information and balanced classes

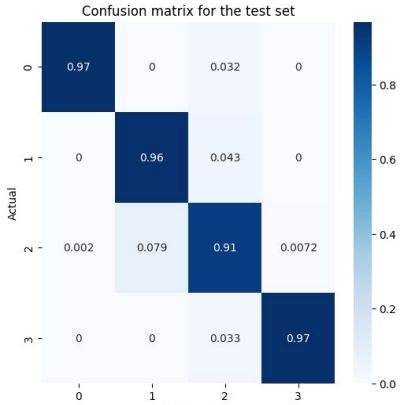


```
"Accuracy":  0.9185667752442996
"Recall":  0.9536846492354968
"F1":  0.8870884573107226
"MCC":  0.787780941145144
"Balanced accuracy":  0.9536846492354968
```

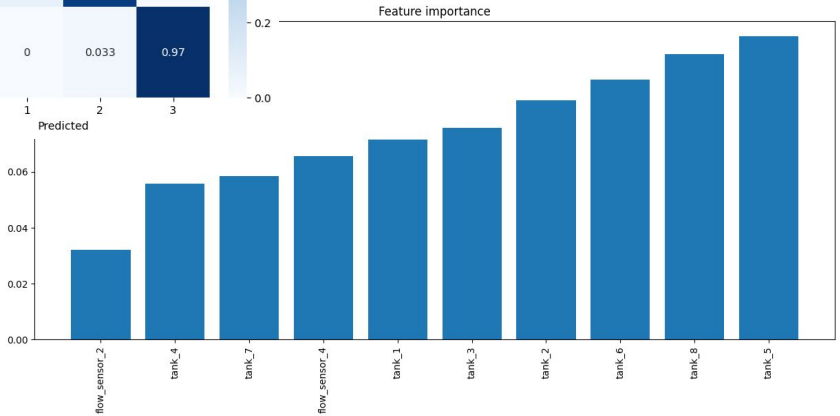

Classifiers - XGBoost

With contextual information and balanced classes

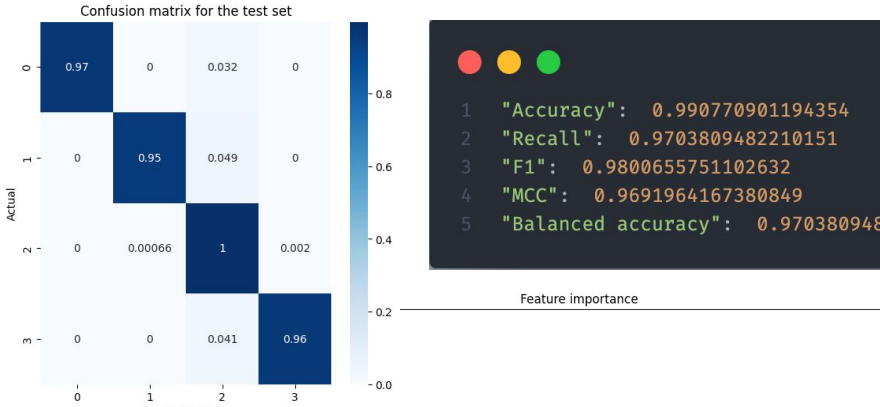


```
"Accuracy":  0.990770901194354
"Recall":  0.9703809482210151
"F1":  0.9800655751102632
"MCC":  0.9691964167380849
"Balanced accuracy":  0.9703809482210151
```

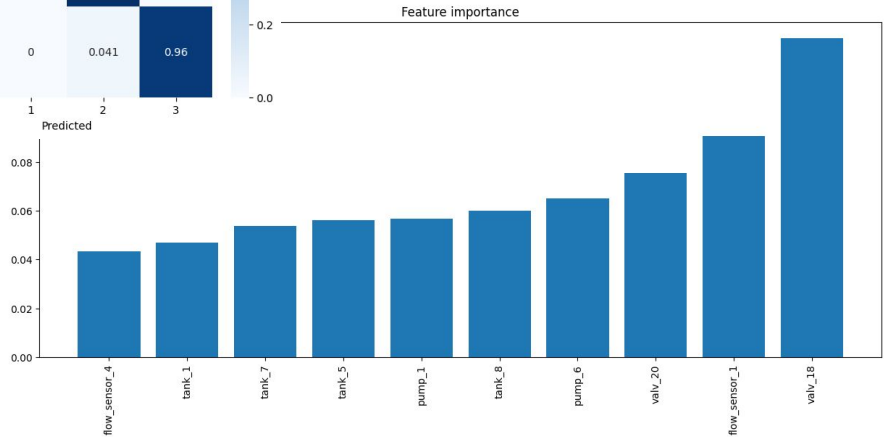

05

Demo

Key results & adversarial attack

06

Conclusion

Recap of results and takeaways

Conclusion

XGBoost is the best

92.3% accuracy on the network test set

Security breach

The network models would still let attacks go through if used in production, so they need to be used carefully

Adversarial Attack

We showed how easy it can be to change the data (and not the model !) and drastically confuse a model => security breach